



Unix  
Bash  
C  
GNU  
Systems

# Software Systems

## Lectures Week 9

### Introduction to Systems Programming 1

(Systems, Concurrency, Inter process communication)

Prof. Joseph Vybihal

Computer Science

McGill University



Unix  
Bash  
C  
GNU  
Systems

## Part A

# Introduction to Systems Programming

COMP 206 – Joseph Vybihal  
Software Systems



# About Systems Programming

Software that interacts with the computer  
and not only with human users

There are three basic ways to interact with  
your system:

- The Shell
- The OS through libraries
- Directly with the devices connected to your machine



# The Shell

There are three basic ways to interact with the shell:

- The command-line arguments (this lecture)
- Executing shell commands (later)
- The shell memory (later)



# C Shell Arguments

Bash-prompt \$ ./a.out 15 Bob 4.2 X

```
int main(int argc, char *argv[]) {  
    :  
    :  
}
```

argc      ← counts the number of arguments  
argv      ← an array of strings, each cell is one argument

In this example:

argc = 4

argv = {"a.out", "15", "Bob", "4.2", "X"}



# C Shell Arguments

Bash-prompt \$ ./a.out 15 Bob 4.2 X

```
int main(int argc, char *argv[]) {  
    int a; char name[30]; float b; char c;  
  
    // Assume I am expecting 4 arguments  
    if (argc != 4) exit(1);  
  
    a = atoi(argv[1]);  
    strcpy(name, argv[2]);    // a string can be copied (or referenced)  
    b = atof(argv[3]);  
    c = *argv[4];             // copies the single character (but really...)  
}
```



# Example

```
#include <stdio.h>
#include <stdlib.h>
int main(int argc, char *argv[]) {
    char c; FILE *p, *q;

    if (argc != 2) exit(1);

    p = fopen(argv[1], "rt");
    q = fopen(argv[2], "wt");
    if (p==NULL || q==NULL) exit(2);

    c = fgetc(p);
    while(!feof(p)) {
        fputc(c, q);
        c = fgetc(p);
    }

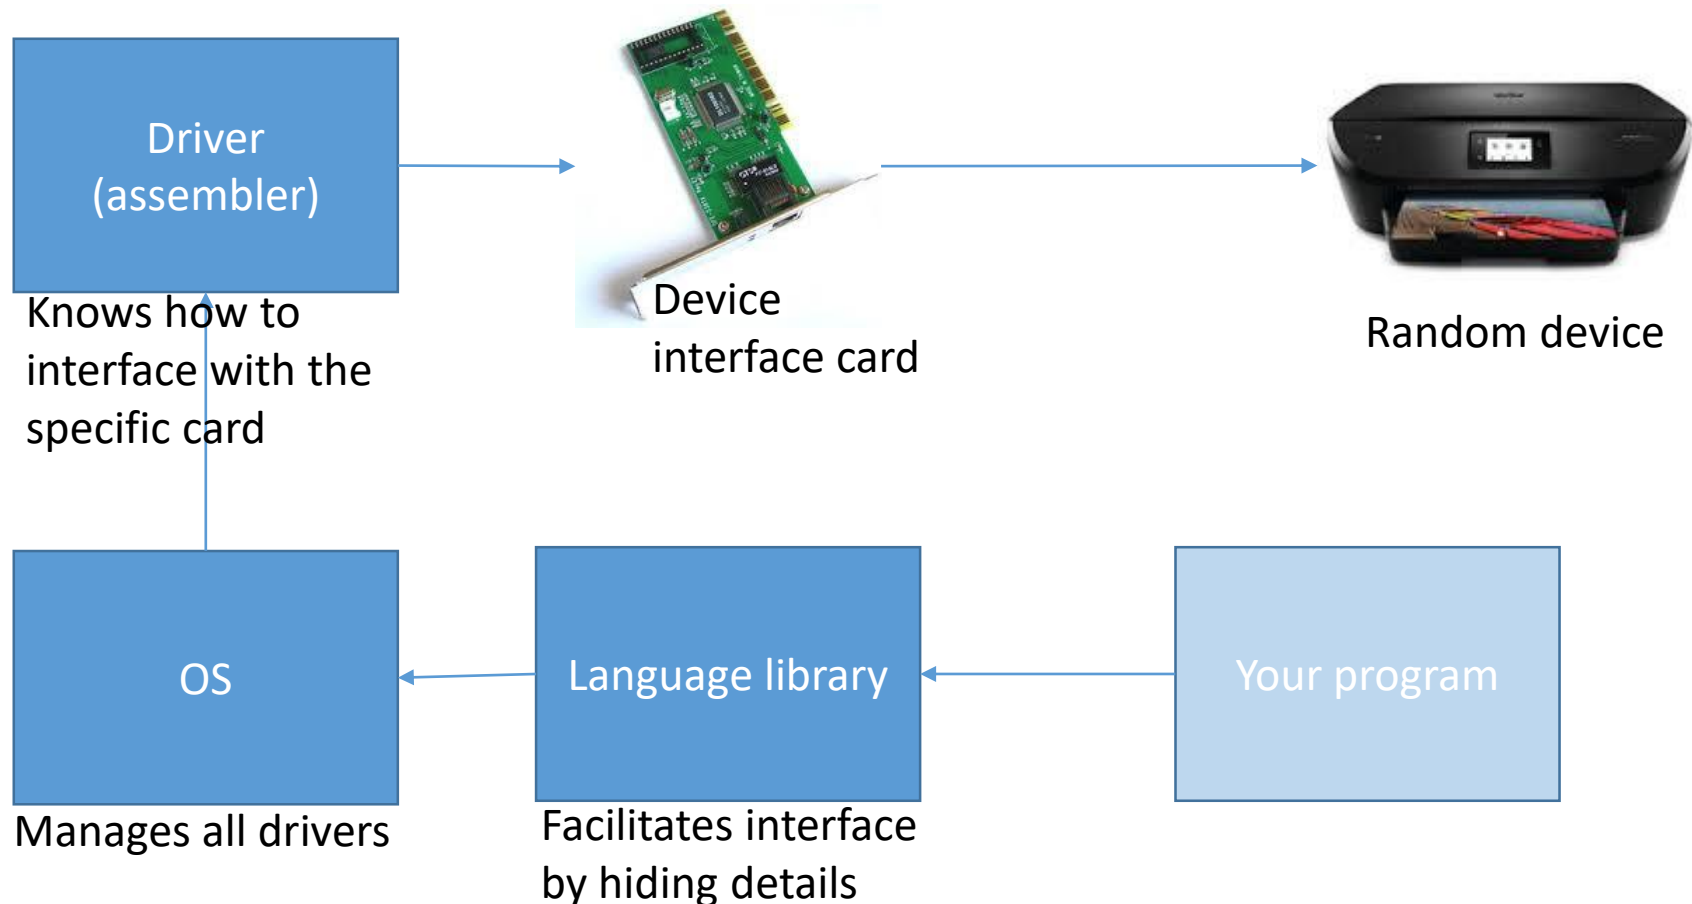
    fclose(p);
    fclose(q);
    return 0;
}
```

```
Bash-prompt $ vi copy.c
Bash-prompt $ gcc -o copy copy.c
Bash-prompt $ ./copy file1.txt file1.bak
```



# Libraries as system interfaces

Every machine connected to your computer passes through the same pipeline:







# Example

```
#include<time.h>
```

- Connects to the system clock
  - `time_t seconds;` // struct stores time since Jan 1 1970
  - `struct tm` // parsed date & time structure
    - `tm_year, tm_mon, tm_mday, tm_hour, tm_min, tm_sec, tm_isdst`
  - `seconds = time(NULL);` // local time since Jan 1 1970
    - `printf("Hours since 1970: %d", seconds/3600);`
  - `float x = difftime(secondsX, secondsY);`
  - `struct tm *t = localtime(&seconds);`
  - `char *p = asctime(t);`
  - `seconds = mktime(t);`



# The time.h library

## struct tm

int	tm_sec	seconds [0,59]
int	tm_min	minutes [0,59]
int	tm_hour	hour [0,23]
int	tm_mday	day of month [1,31]
int	tm_mon	month of year [0,11]
int	tm_year	years since 1900
int	tm_wday	day of week [0,6] (Sunday = 0)
int	tm_yday	day of year [0,365]
int	tm_isdst	daylight savings flag

## time\_t

An unsigned long integer or unsigned long double number (depends on implementation) that measures the number of milliseconds from a fixed point in time to the present as an offset (or distance measurement).

Jan 1 1970 is the earliest date/time that can be represented. Jan 1 1970 = 0.



# The time.h library

```
char *    asctime(const struct tm *);
char *    asctime_r(const struct tm *, char *);
clock_t   clock(void);
int        clock_getres(clockid_t, struct timespec *);
int        clock_gettime(clockid_t, struct timespec *);
int        clock_settime(clockid_t, const struct timespec *);
char *    ctime(const time_t *);
char *    ctime_r(const time_t *, char *);
double     difftime(time_t, time_t);
struct tm *getdate(const char *);
struct tm *gmtime(const time_t *);
struct tm *gmtime_r(const time_t *, struct tm *);
struct tm *localtime(const time_t *);
struct tm *localtime_r(const time_t *, struct tm *);
time_t     mktime(struct tm *);
int        nanosleep(const struct timespec *, struct timespec *);
size_t     strftime(char *, size_t, const char *, const struct tm *);
char *     strptime(const char *, const char *, struct tm *);
time_t     time(time_t *);
int        timer_create(clockid_t, struct sigevent *, timer_t *);
int        timer_delete(timer_t);
int        timer_gettime(timer_t, struct itimerspec *);
int        timer_getoverrun(timer_t);
int        timer_settime(timer_t, int, const struct itimerspec *, struct itimerspec *);
void       tzset(void);
```



# Example

## Using time.h as a profiler.

```
#include <stdio.h>
#include <time.h>
int main()
{
    time_t begin,end;
    long i;

    begin= time(NULL);
    for(i = 0; i < 150000000; i++);
    end = time(NULL);

    printf("for loop used %f seconds to complete the execution\n", difftime(end, begin));

    return 0;
}
```



# Directly with devices connected to the computer



As a system's language, C permits direct access to devices.

- A device card has an address
  - `void *p = 145; // assume we know the card's address is 145`
- We can communicate with the device
  - `*p = 'A'; // if p points to printer then printer prints 'A'`
  - `int x = *p; // if p points to printer status then x = status`
- Devices speak in binary...



# Binary

## Counting in decimal and binary:

0	00000000
1	00000001
2	00000010
3	00000011
4	00000100
5	00000101

In Decimal: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9

In Binary : 0, 1

In decimal:  $9 + 1 = 10$  (we reuse digits, = 10)

In binary:  $1 + 1 = 10$  (we reuse digits, = 2)

Binary is used in many ways:

- The char is ASCII which is coded binary
- Numbers are stored as binary
- Binary 00101 can be thought of as three false values and two true values



# Manipulating binary in C

## Operators:

- & Binary and
- | Binary or
- ~ Binary complement
- >> Binary shift right
- << Binary shift left



# Manipulating binary in C

Meaning:

$\&$       $1011 \& 0110 \rightarrow 0010$

$|$       $1011 | 0110 \rightarrow 1111$

$\sim$       $\sim 1011 \rightarrow 0100$

$\gg$       $1011 \gg 3 \rightarrow 0001$

$\ll$       $1011 \ll 2 \rightarrow 1100$

AND:

A	B	AND	R
1	1		1
1	0		0
0	1		0
0	0		0

OR:

A	B	OR	R
1	1		1
1	0		1
0	1		1
0	0		0

Complement:  
Means oposite





# Manipulating binary in C

Usage:

```
int a = 11; // Binary 1011
```

```
int b = 6;  // Binary 0110
```

```
int r;
```

```
r = a & b; // 1011 & 0110 → 0010
```

```
r = a | b; // 1011 | 0110 → 1111
```

```
r = ~a;    // ~1011 → 0100
```

```
r = a >> 3; // 1011 >> 3 → 0001
```

```
r = a << 2; // 1011 << 2 → 1100
```



# Example

## Masking:

- When we want to change a single bit, or
- When we want to find out about a bit.

```
int b = 6; // Binary 0110
```

```
b = b | 0001; // "set" the last bit to 1
```

```
b = b & 1110; // "set" the last bit to 0
```

```
if (b & 0001) // =0 if last bit was 0, else =1
```



Unix  
Bash  
C  
GNU  
Systems

## Part B

# Concurrent Programs



# What is a concurrent program?

It is when two programs are running at the same time on the same machine.

It does not mean that they are cooperating, but often they are.



# Ampersand and semi-colon

- Sequential command-line execution
  - Example:
    - `ls; cp file1.c /backup; cat file1.c`
  - In the above example the `ls` command is executed first. After the `ls` command is finished the `cp` command executes. After `cp` is finished then the `cat` command runs last.
- Concurrent command-line execution
  - Example:
    - `ls & cp file1.c /backup & cat file1.c`
  - In this example all three commands execute concurrently.
  - They do not launch at the same time – if they run long enough they will all use the CPU at the same time.
  - Any concurrent output is displayed to the screen at the same time... mixed together.



# Using the Shell for Concurrency

- ;
  - Sequential execution of commands
  - Example: `who; grep 'Jack' eg.txt; ls > out.txt`
- &
  - Parallel execution of commands
  - Example: `who & whoami & ls`

What would the output look like when invoked on the command line?



# Related Commands

- **PS**

- See all the currently running programs (processes). Notice that each process has an ID number called the PID.
- Syntax to see your own: `ps`
- Syntax to see everyone: `ps -e`

- **KILL**

- Terminate an executing program (process)
- Syntax standard : `kill PID`
- Syntax emergency: `kill -9 PID`



# Active Processes

- The `ps` command is an ideal solution for troubleshooting problem processes.
- Although the command options have a tendency to change from one OS to another, here are some of the common options.
  - `-a` : all processes, all users
  - `-e` : environment/everything
  - `-g` : process group leaders as well
  - `-l` : long format
  - `-u` : user oriented report
  - `-x` : even processes not executed from terminals
  - `-f` : full listing





# Demo

- Semi-colon
- Ampersand
- PS
- Kill



# Example

Bash-prompt \$ vi producer.c

Bash-prompt \$ vi consumer.c

Bash-prompt \$ gcc -o producer producer.c

Bash-prompt \$ gcc -o consumer consumer.c

Bash-prompt \$ ./producer & ./consumer

These two programs runs concurrently without using fork().

These use a text file to coordinate and share the data (as we saw last class).



# Example

```
// PRODUCER.C
int main() {
    char c = ' ';
    FILE *p;

    while (c != 'x') {
        c = getchar();

        while ((p=fopen("shared.txt","at") == NULL) ;
            fprintf(p,"%c\n", c);
            fclose(p);
        }

    return 0;
}
```



```
// CONSUMER.C
int main() {
    char c; int pos = 0;
    FILE *q;

    do {
        c = removeCharacter(pos);
        printf("%c", c);
        pos++;
    } while (c != 'x');
}

char removeCharacter(int pos) {
    char c;
    FILE *q;
    int x = 0;

    while ((q=fopen("shared.txt", "rt")) == NULL) ;

    c=fgetc(q);
    while(!feof(q) && x < pos) {
        pos++;
        c=fgetc(q);
    }
    fclose(q);
    return c;
}
```



# Accessing Shell Memory

Remember:

Bash-prompt \$ set

The 'set' command displays all the variables defined within the shell memory.

C can access these variables:

```
#include <stdlib.h>
char *data = getenv("VARIABLE_NAME");
```

```
printf("PATH : %s\n", getenv("PATH"));
printf("HOME : %s\n", getenv("HOME"));
printf("ROOT : %s\n", getenv("ROOT"));
```

The getenv() returns NULL if it failed to find the variable.

```
int n = atoi(getenv("CONTENT_LENGTH"));
```

```
int setenv (const char *name, const char *value, int replace)
```

```
int success = setenv("VAR_NAME", "VALUE", 1);
```

Replace = 1 for true, 0 for false  
Success = 0 for true, -1 for false



# Example

Bash-prompt `$ ./prog1; ./prog2`

Notice that prog1 runs to completion first before prog2 starts.

In prog1:

- `setenv("x", "yes", 1);`

In prog2:

- `char *p = getenv("x");`

Notice that a preceding program can communicate with a following program.



# Example

Bash-prompt `$ set x="yes"`

User sets the variable at the command prompt

Bash-prompt `$ ./prog2`

In prog2:

- `char *p = getenv("x");`

Notice that the user can change the “environment” causing prog2 to behave differently.



# Example

Bash-prompt `$ ./prog3 & ./prog4`

Two concurrent programs can coordinate using the shell memory

Assume a shell variable `TURN` is “3” or “4”.

In prog3:

```
char *p;
while (1) {
    do { p = getenv("TURN"); }
    while (strcmp(p, "3") != 0) ; // wait
    ... do something since it is 3 now ...
    setenv("TURN", "4", 1);
}
```

In prog4:

```
char *q;
while (1) {
    do { q = getenv("TURN"); }
    while (strcmp(q, "4") != 0) ; // wait
    ... do something since it is 4 now ...
    setenv("TURN", "3", 1);
}
```





# Question

Suggest a way we could use shell memory to do the producer/consumer problem. Use the ampersand to launch the two programs.

In this case we would like to simply produce one value and consume that one value, and repeat. The user enters values from the keyboard until they write the word DONE, which then terminates both programs.



Unix  
Bash  
C  
GNU  
Systems

## Part C

# Inter-process Communication

Readings: <http://www.csl.mtu.edu/cs4411.ck/www/NOTES/process/fork/create.html>



# What is a process?

## Program

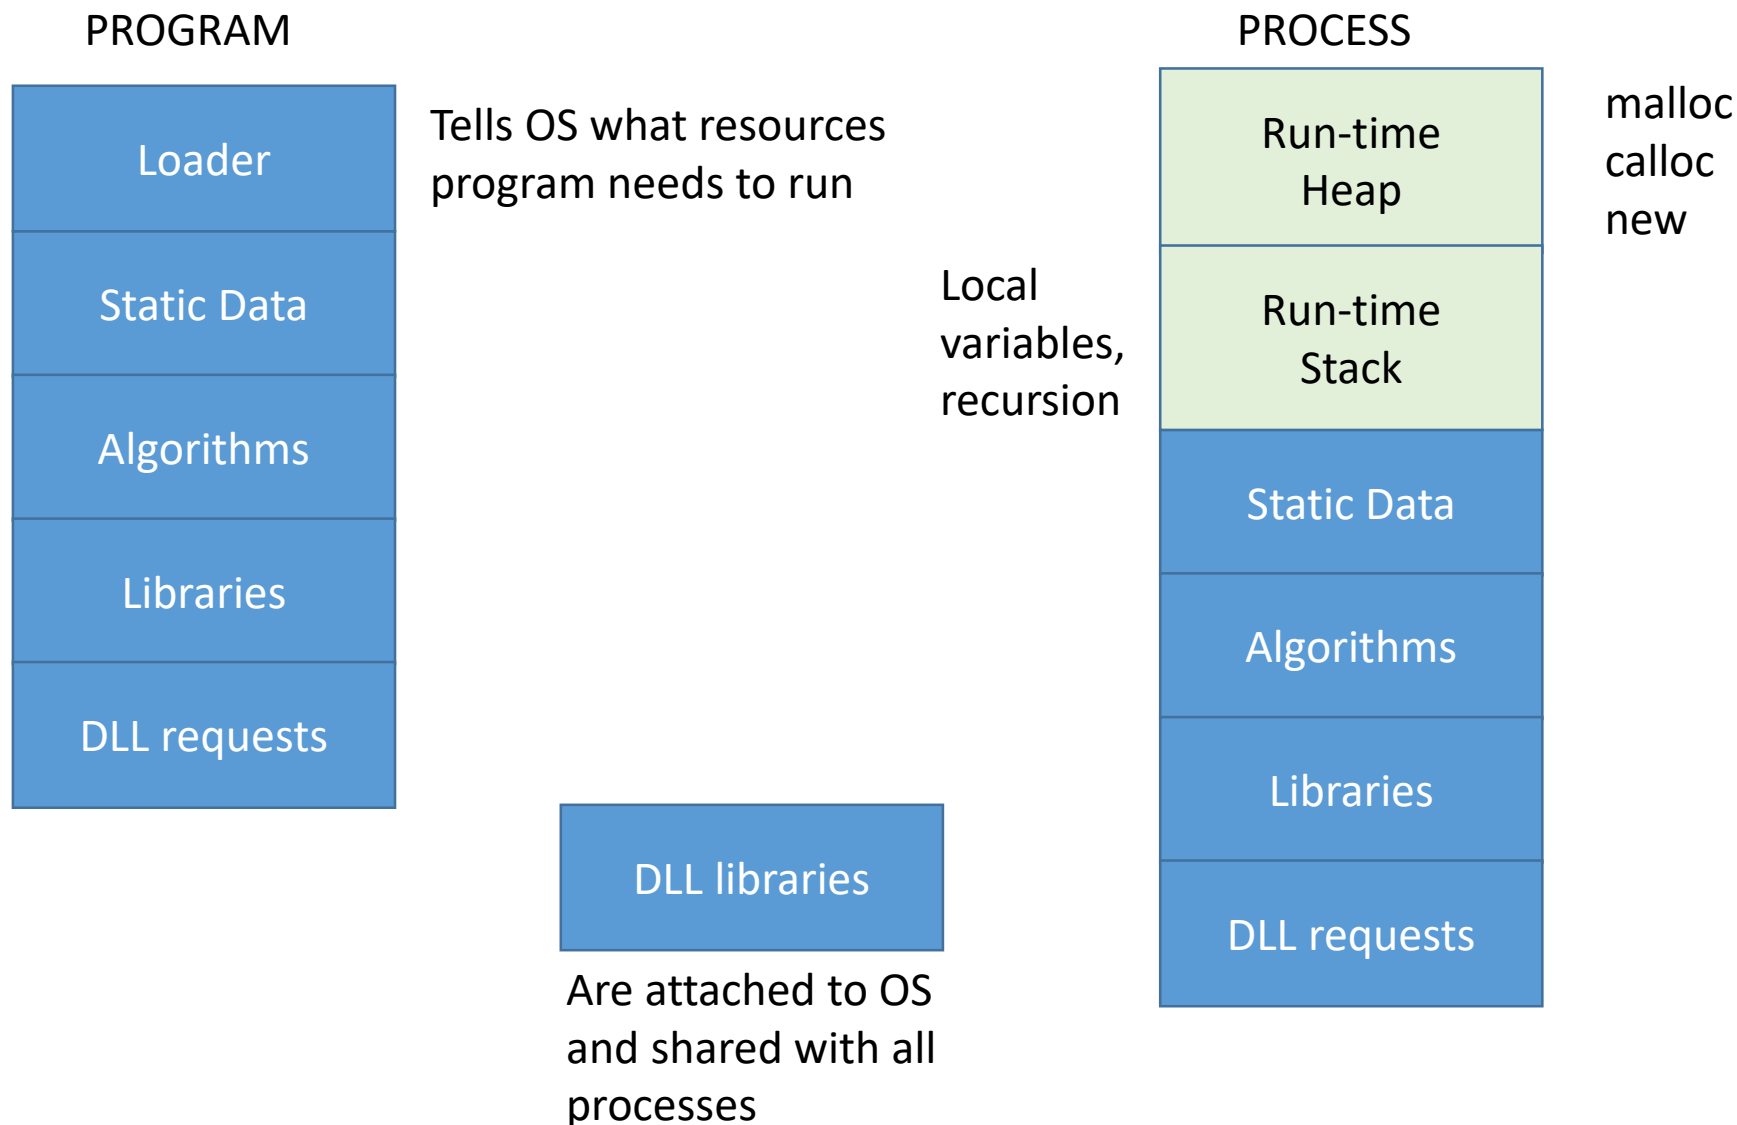
- A program is understood to mean a file on disk containing a runnable algorithm (compiled or interpret-able)

## Process

- A copy of the program but in RAM (the computer's live memory) and it is currently being executed



# Program & Process





# Three types of processes

- **Shell processes**
  - Is a process that runs within its own shell. Launching a process of this form will pause the current shell by launching a new shell and then using the command-line of the new shell to launch a program or execute a shell command.
- **Cloned process**
  - The currently executing process can clone itself. After the clone operation there are two identical processes running in the same shell concurrently.
- **New process**
  - The currently executing process can launch a new program within the current shell. Both processes run concurrently.

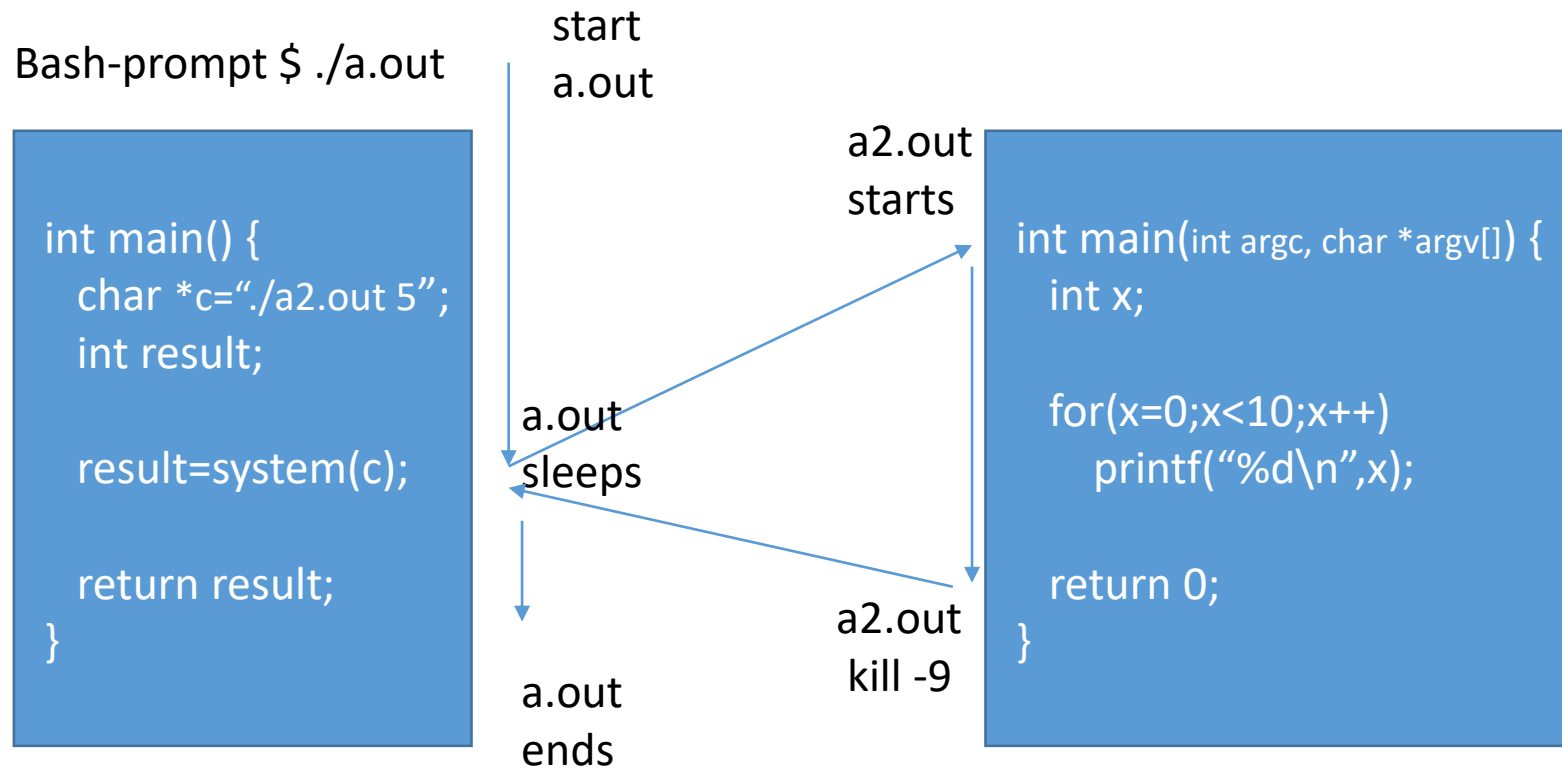


# Three types of processes

- Shell processes
  - `#include<stdlib.h>`
  - `int system(char *command_line_command);`
  - Returns -1 on error, or return status number
- Cloned process
  - `#include<unistd.h>`
  - `int fork();`
  - Returns the process ID number
- New process
  - Not covered...



# The system() function



Programs can use system() as often as needed.

A process invoked by system() can also use system(), recursive definition.

Notice how data is passed through command-line arguments & return statement.



# The fork() function

```
#include <stdlib.h> /* needed to define exit() */
#include <unistd.h> /* needed for fork() and getpid() */
#include <stdio.h> /* needed for printf() */

int main(int argc, char *argv[]) {
    int pid; /* process ID */

    switch (pid = fork()) {
        case 0: /* fork returns 0 to the child */
            printf("I am the child process: pid=%d\n", getpid());
            break;

        default: /* fork returns the child's pid to the parent */
            printf("I am the parent process: pid=%d, child pid=%d\n", getpid(), pid);
            break;

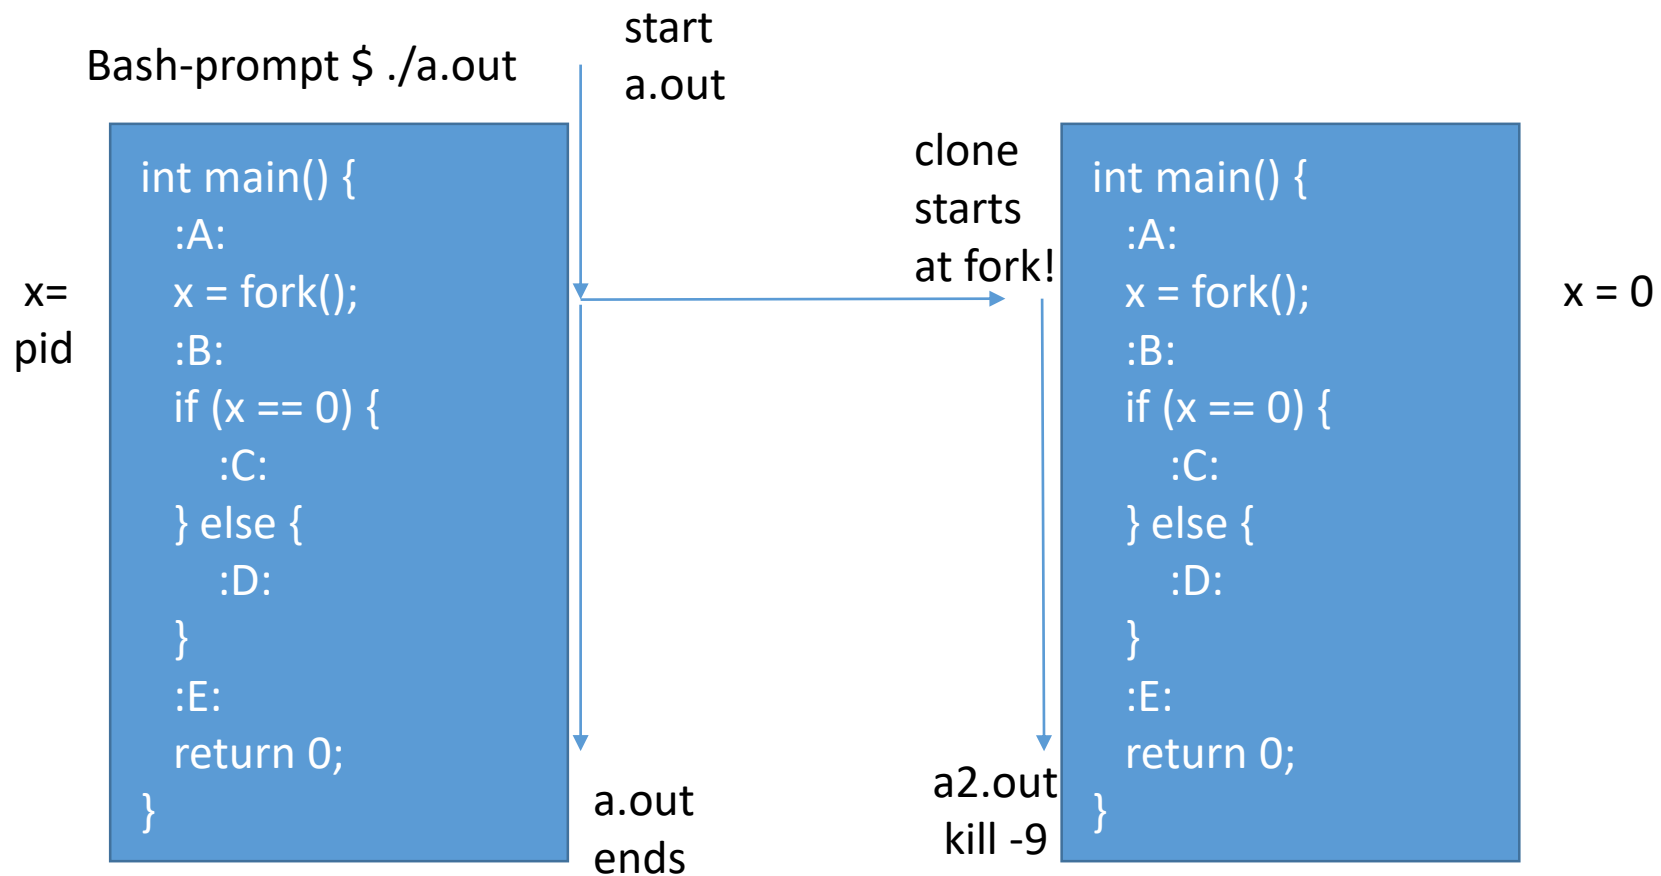
        case -1: /* something went wrong */
            perror("fork"); // send string to STDERR
            exit(1);
    }
    exit(0);
}
```

Programs can use fork() as often as needed.  
A process invoked by fork() can also use fork().





# The fork() function



a.out → :A: :B: :D: :E:

clone → :B: :C: :E:

The clone receives a copy of all the variables (it is not shared).



# The producer & consumer problem

## The producer

- A program that generate data saving that data in a data structure or in a file.

## The consumer

- A program that reads data from a data structure or a file performing some kind of computation.

## Concurrency

- For some reason the producer and consumer need to run independently but cooperatively.
  - Maybe the computation is slow but the data comes in quickly. The intermediate data structure or file is used as a temporary cache.



# Example

```
#include <stdlib.h> #include <unistd.h> #include <stdio.h>
```

```
int main() {  
    int pid = fork();  
    if (pid == -1) exit(1);  
    if (pid == 0) { producer(); wait(); } // wait for consumer to end  
    if (pid != 0) { consumer(); wait(); } // wait for child to end  
}
```

```
void producer() {  
    :  
    :  
}
```

```
void consumer() {  
    :  
    :  
}
```

Not an efficient example, but easy to understand.

Notice the use of functions to make the code easier to read.



# Example

```
void producer() {
    char c = ' ';
    FILE *p;

    while (c != 'x') {                // program stops at the input of 'x'
        c = getchar();

        while ((p=fopen("shared.txt","at")) == NULL) ; // notice semi-colon, we wait...
        fprintf(p,"%c\n", c);
        fclose(p);                    // give the other process a change to open the file
    }
}

void consumer() {
    char c; int pos = 0;
    FILE *q;

    do {
        c = removeCharacter(pos); // we encapsulate further here for understandability
        printf("%c", c);
        pos++;                      // pos tracks the next data, pos is independent of producer
    } while (c != 'x');             // program stops at input of 'x'
}
```

Data sharing is handled  
by a text file cache.



# Example

How can we do  
this with fseek?

```
char removeCharacter(int pos) {  
    char c;  
    FILE *q;  
    int x = 0;  
  
    while ((q=fopen("shared.txt", "rt")) == NULL) ; // busy wait  
  
    c=fgetc(q);  
    while(!feof(q) && x < pos) {  
        x++;                                     // move to the correct position in file  
        c=fgetc(q);  
    }  
  
    fclose(q);  
  
    return c;                                   // return the character  
}
```

Notice that the two programs run independently from each other.  
The speed by which they processed the file was different and did not matter.  
However the current implementation assumes that producer is faster than consumer.



# Question

The previous example assumed that producer was faster than consumer.

How could we modify consumer to handle the case of being sometimes faster than producer?



# Using shared memory

```
#include <sys/types.h>
```

```
#include <sys/shm.h>
```

```
void *shmat(int shmid, const void *shmaddr, int shmflg);
```

- Attach shared memory to program

```
int shmdt(const void *shmaddr);
```

- Detach shared memory from program

```
int shm_id = shmget(  
    key_t    k,      /* the key for the segment */  
    int     size,    /* the size of the segment */  
    int     flag);   /* create/use flag */
```

- Creates the shared memory space



# Example

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include <stdio.h>
```

```
.....
```

```
int    shm_id;    /* shared memory ID    */
struct MEM { int a, b, c, status; } *p; // status=0 empty, 1 filled
```

```
.....
```

```
shm_id = shmget(IPC_PRIVATE, sizeof(struct MEM), IPC_CREAT | 0666);
if (shm_id < 0) exit(1);
```

```
/* now the shared memory ID is stored in shm_id */
```

```
p = (int *) shmat(shm_id, NULL, 0);
if (p == -1) exit(1);
```

```
p->a = 5;
p->status = 1;
while(p->status == 1) sleep(1); // wait for child, child will change status to 0
exit(0);
```

**IPC\_CREAT | 0666** for a server (*i.e.*, creating and granting read and write access to the server) - Octal

**0666** for any client (*i.e.*, granting read and write access to the client) - Octal

If a client wants to use a shared memory created with **IPC\_PRIVATE**, it must be a child process

Child program is similar.  
Examples: [HERE](#) and [HERE](#).